



Akri

Presented by Kate Goldenring and Jiri Appl

Who We Are



Jiri Appl

Principal Software
Engineer Microsoft



@jiriappl



@jiria



Kate Goldenring

Software Engineer
Microsoft




@kate-goldenring



What can run Kubernetes on the edge?

1. Devices that have enough compute power to run Kubernetes effectively
2. Devices that are expected to be used as general compute

Heavy Edge



Server Class



Light Edge



Gateway Class
Industrial PCs



Tiny Edge




MCU Class
Smart Camera

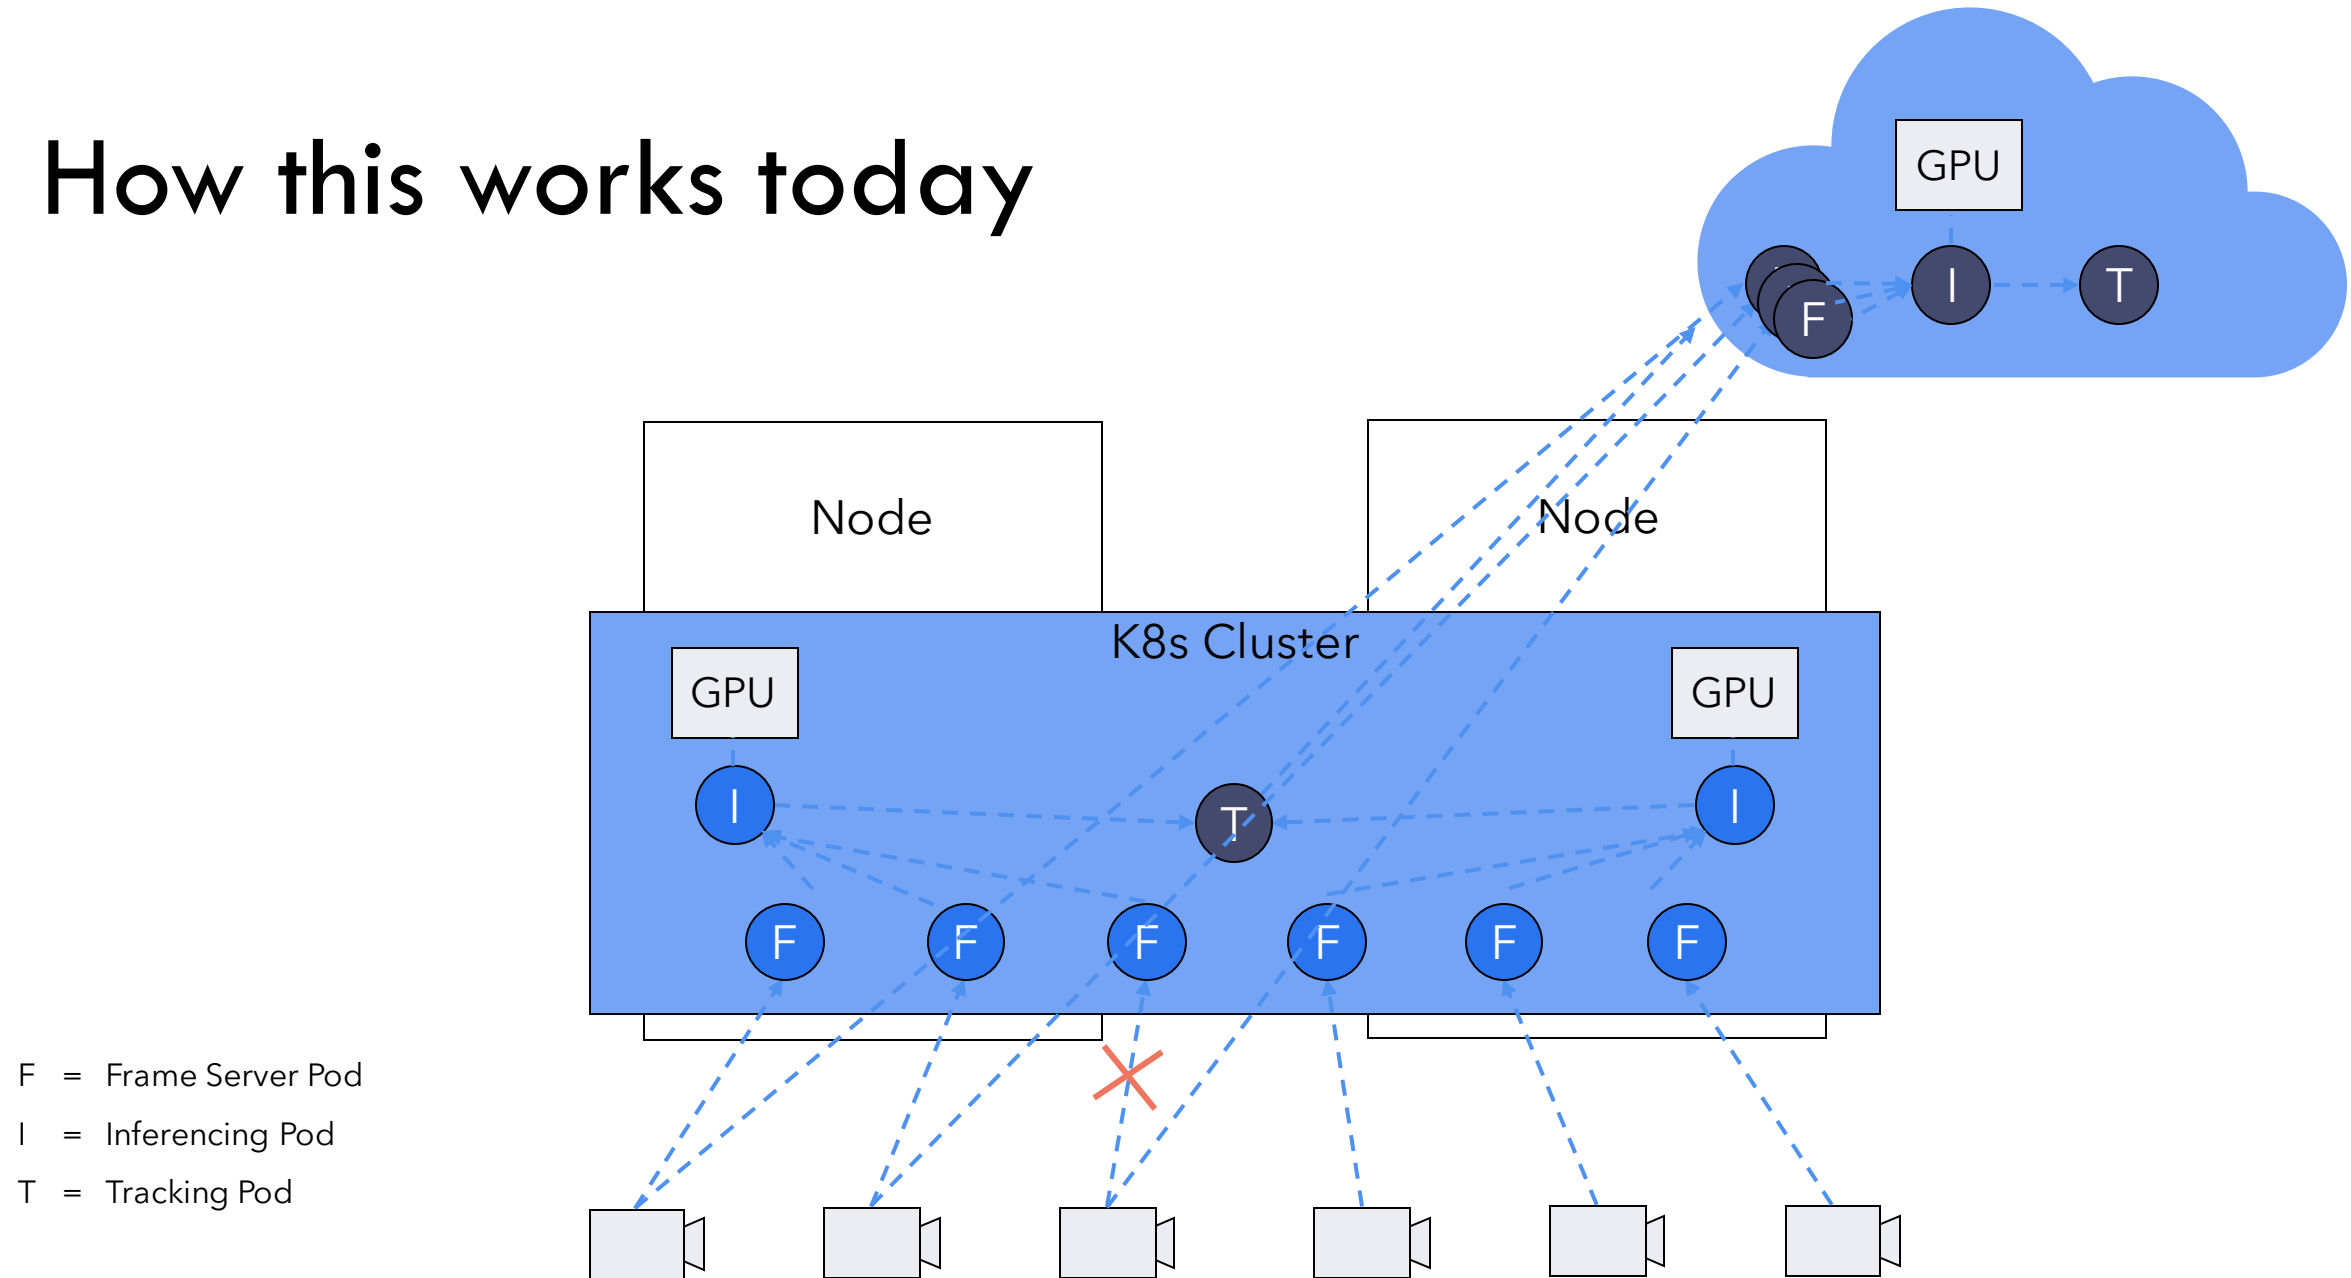


Scenario Example

- Smart Retail
 - Existing Kubernetes Cluster, GPU capable nodes
 - IP based Cameras on the network (ONVIF)
- Kubernetes Application for Social Distancing
 - How can I direct traffic throughout my retail store to ensure social distancing?



How this works today



Akri – A Kubernetes Resource Interface (for the Edge)

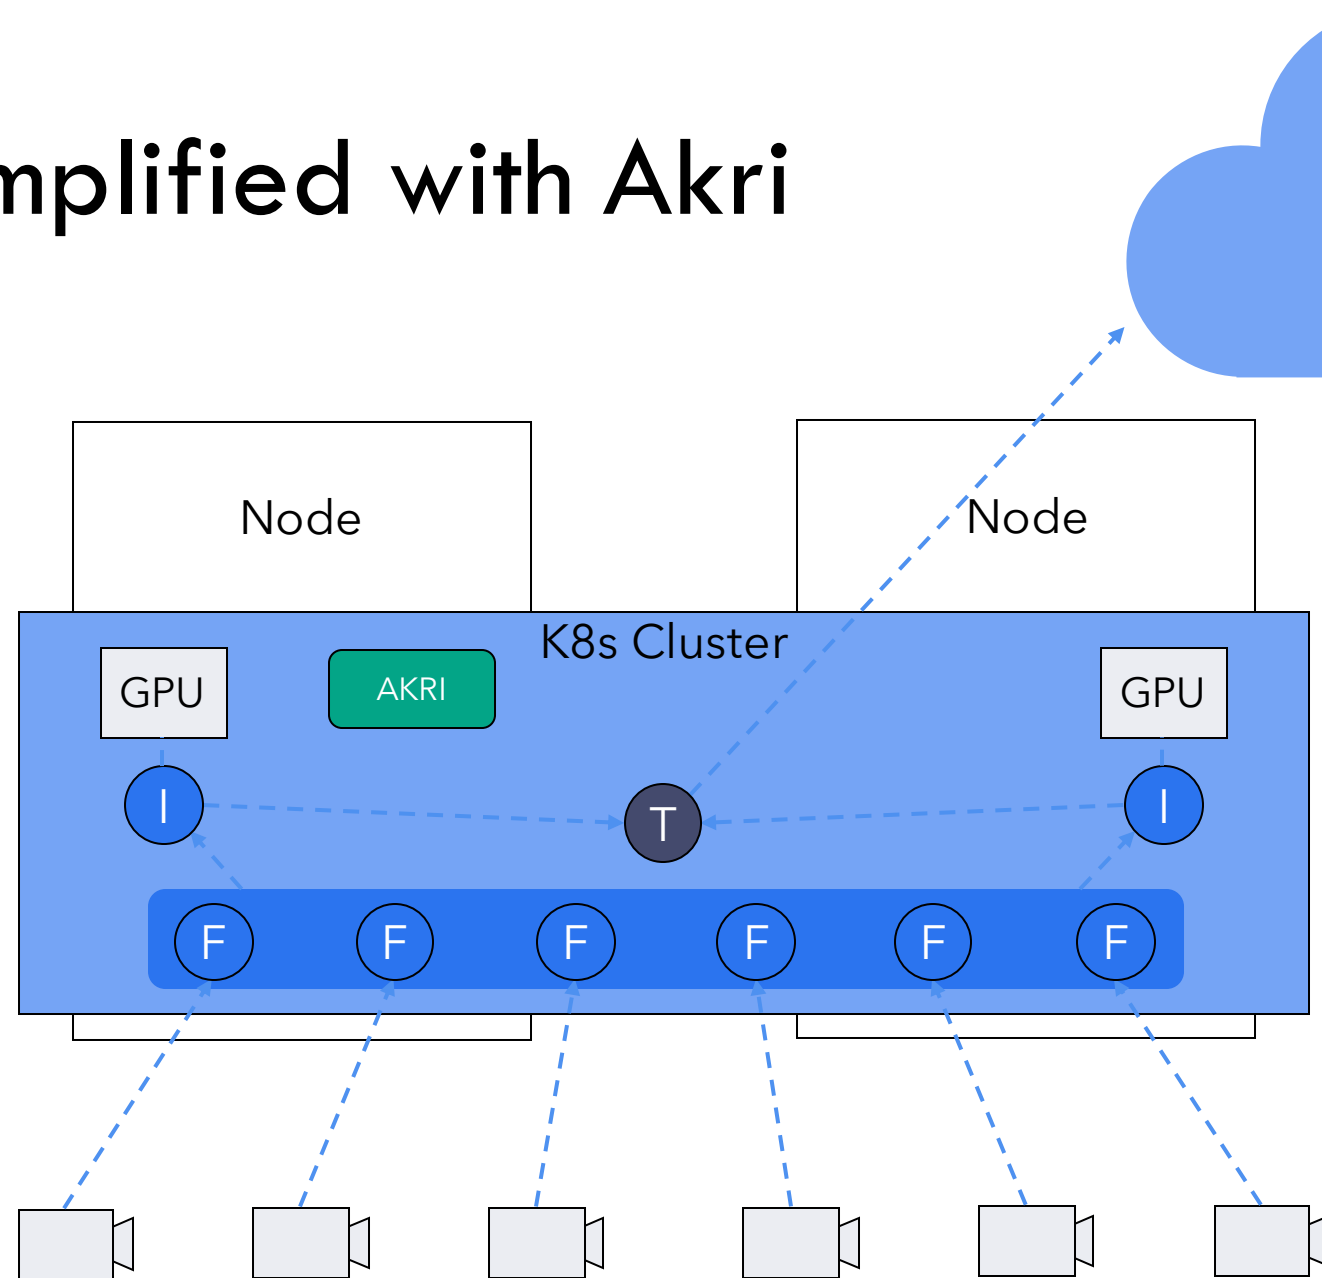
- Open source project that attempts to define a Kubernetes based approach on representing leaf devices (such as IP cameras and USB devices) found on the edge as native Kubernetes resources
- It also supports the exposure of simple embedded hardware resources
- Schedules workloads based on the availability of leaf devices
- Utilizes Kubernetes device plugin framework
- Edge optimized for low footprint clusters-
Implemented in Rust



Scenario simplified with Akri

```
kind: Configuration
metadata:
  name: akri-udev-gpu
spec:
  protocol:
    udev:
  brokerPodSpec:
    containers:
      - name: akri-udev-gpu-broker
        image: "ghcr.io/.../ml-inference"
```

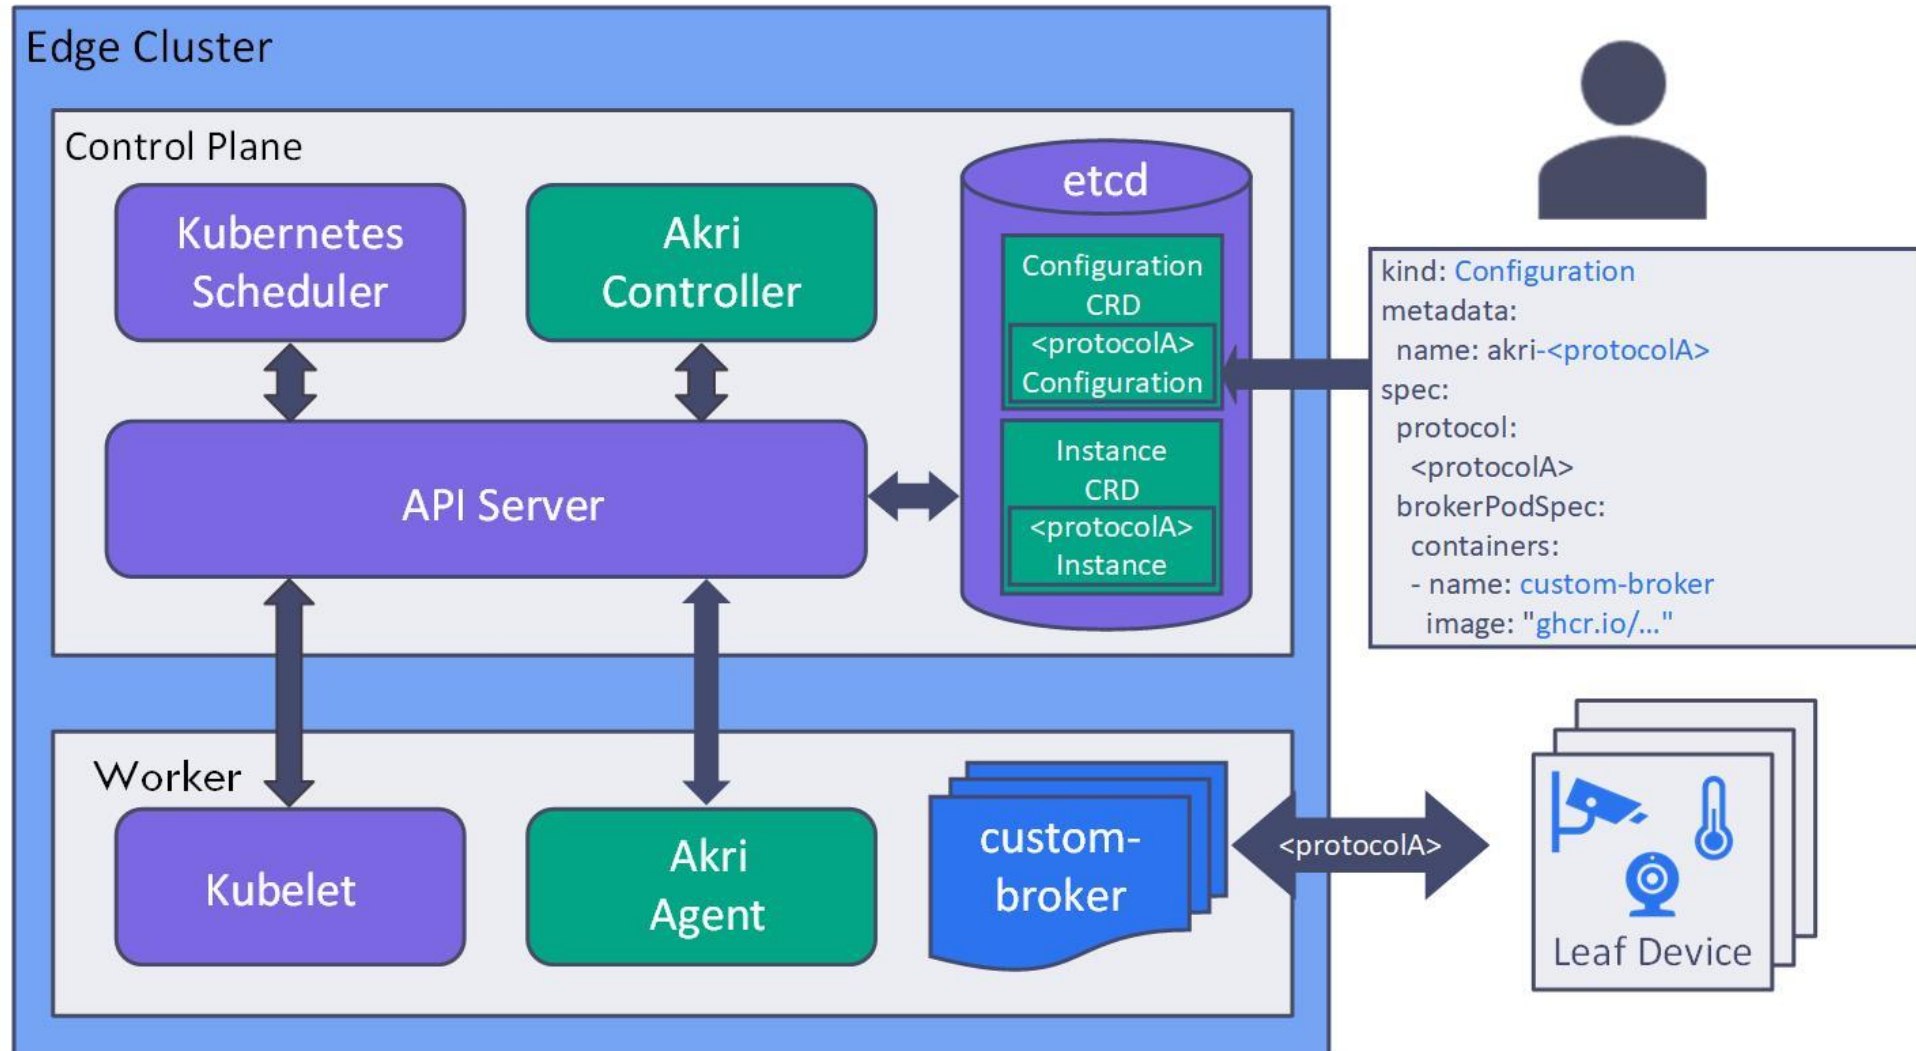
F = Frame Server Pod
I = Inferencing Pod
T = Tracking Pod



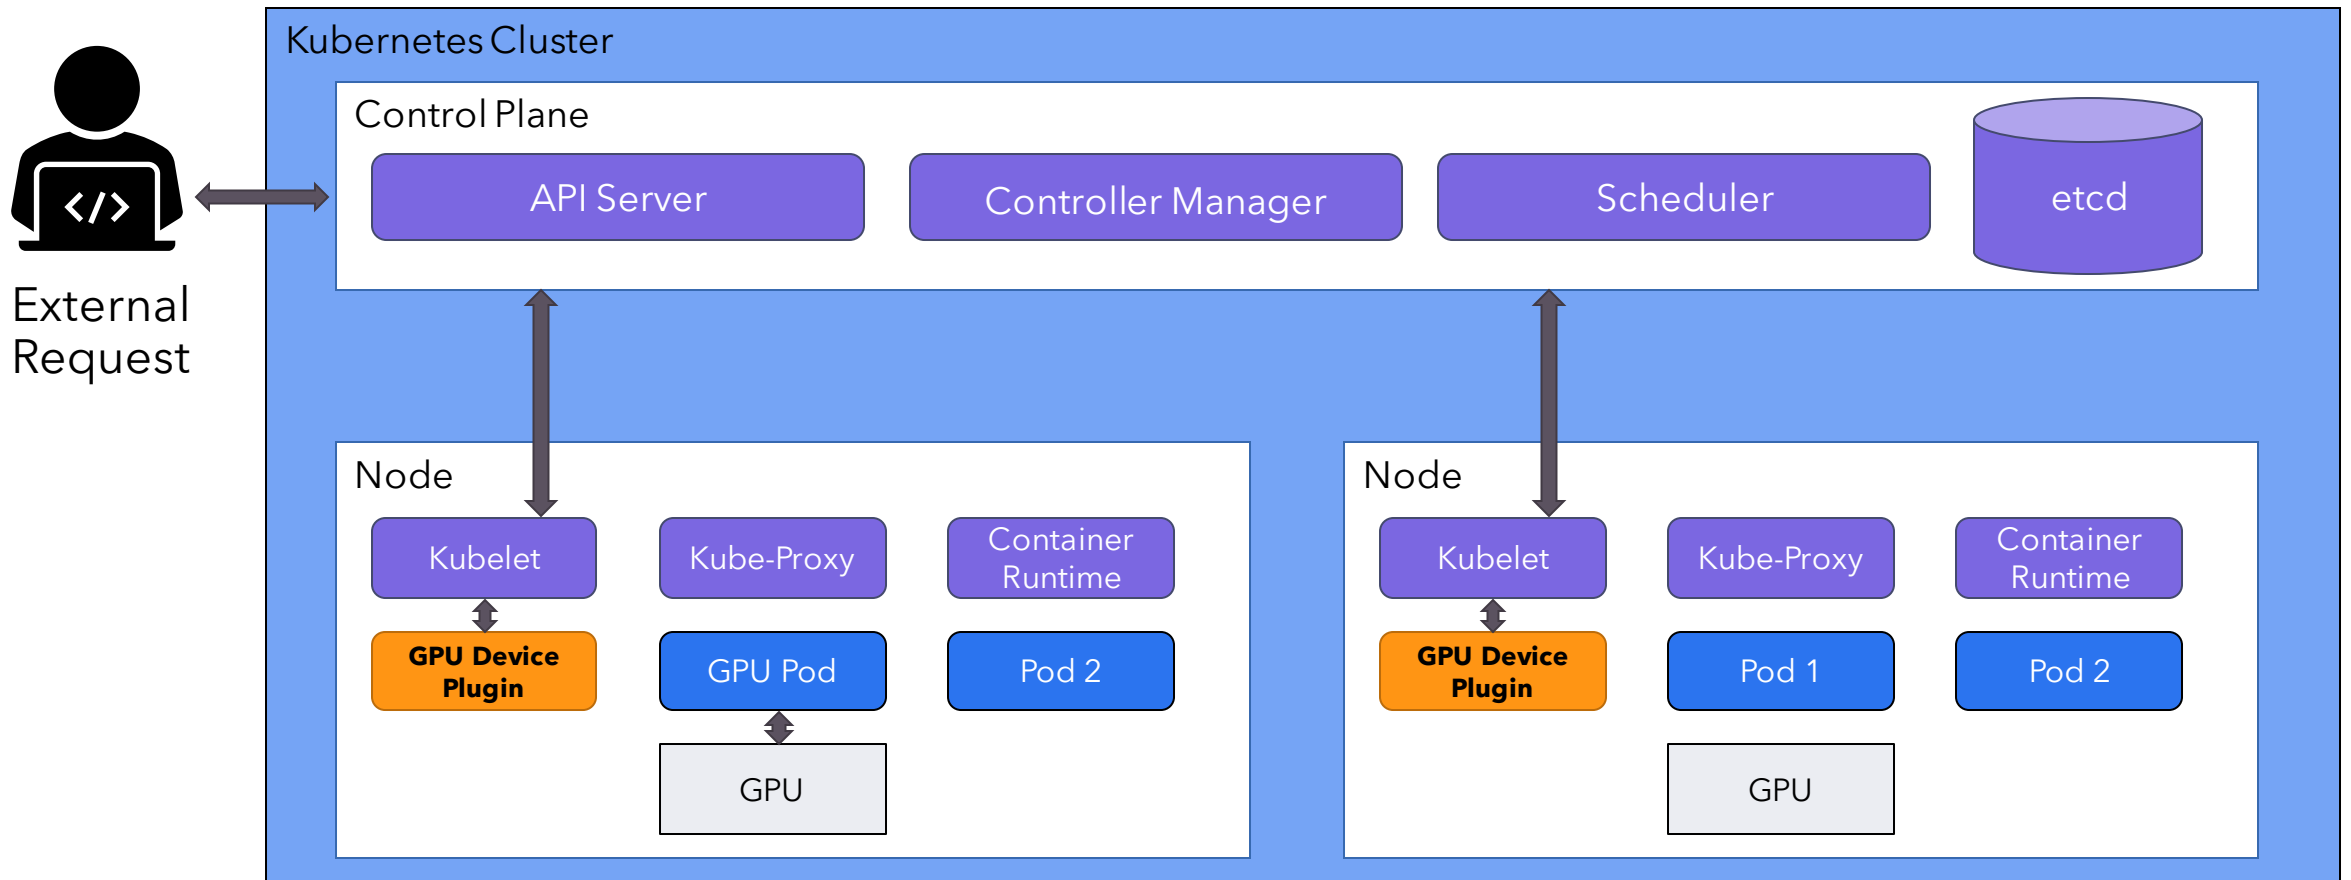


Technical Overview

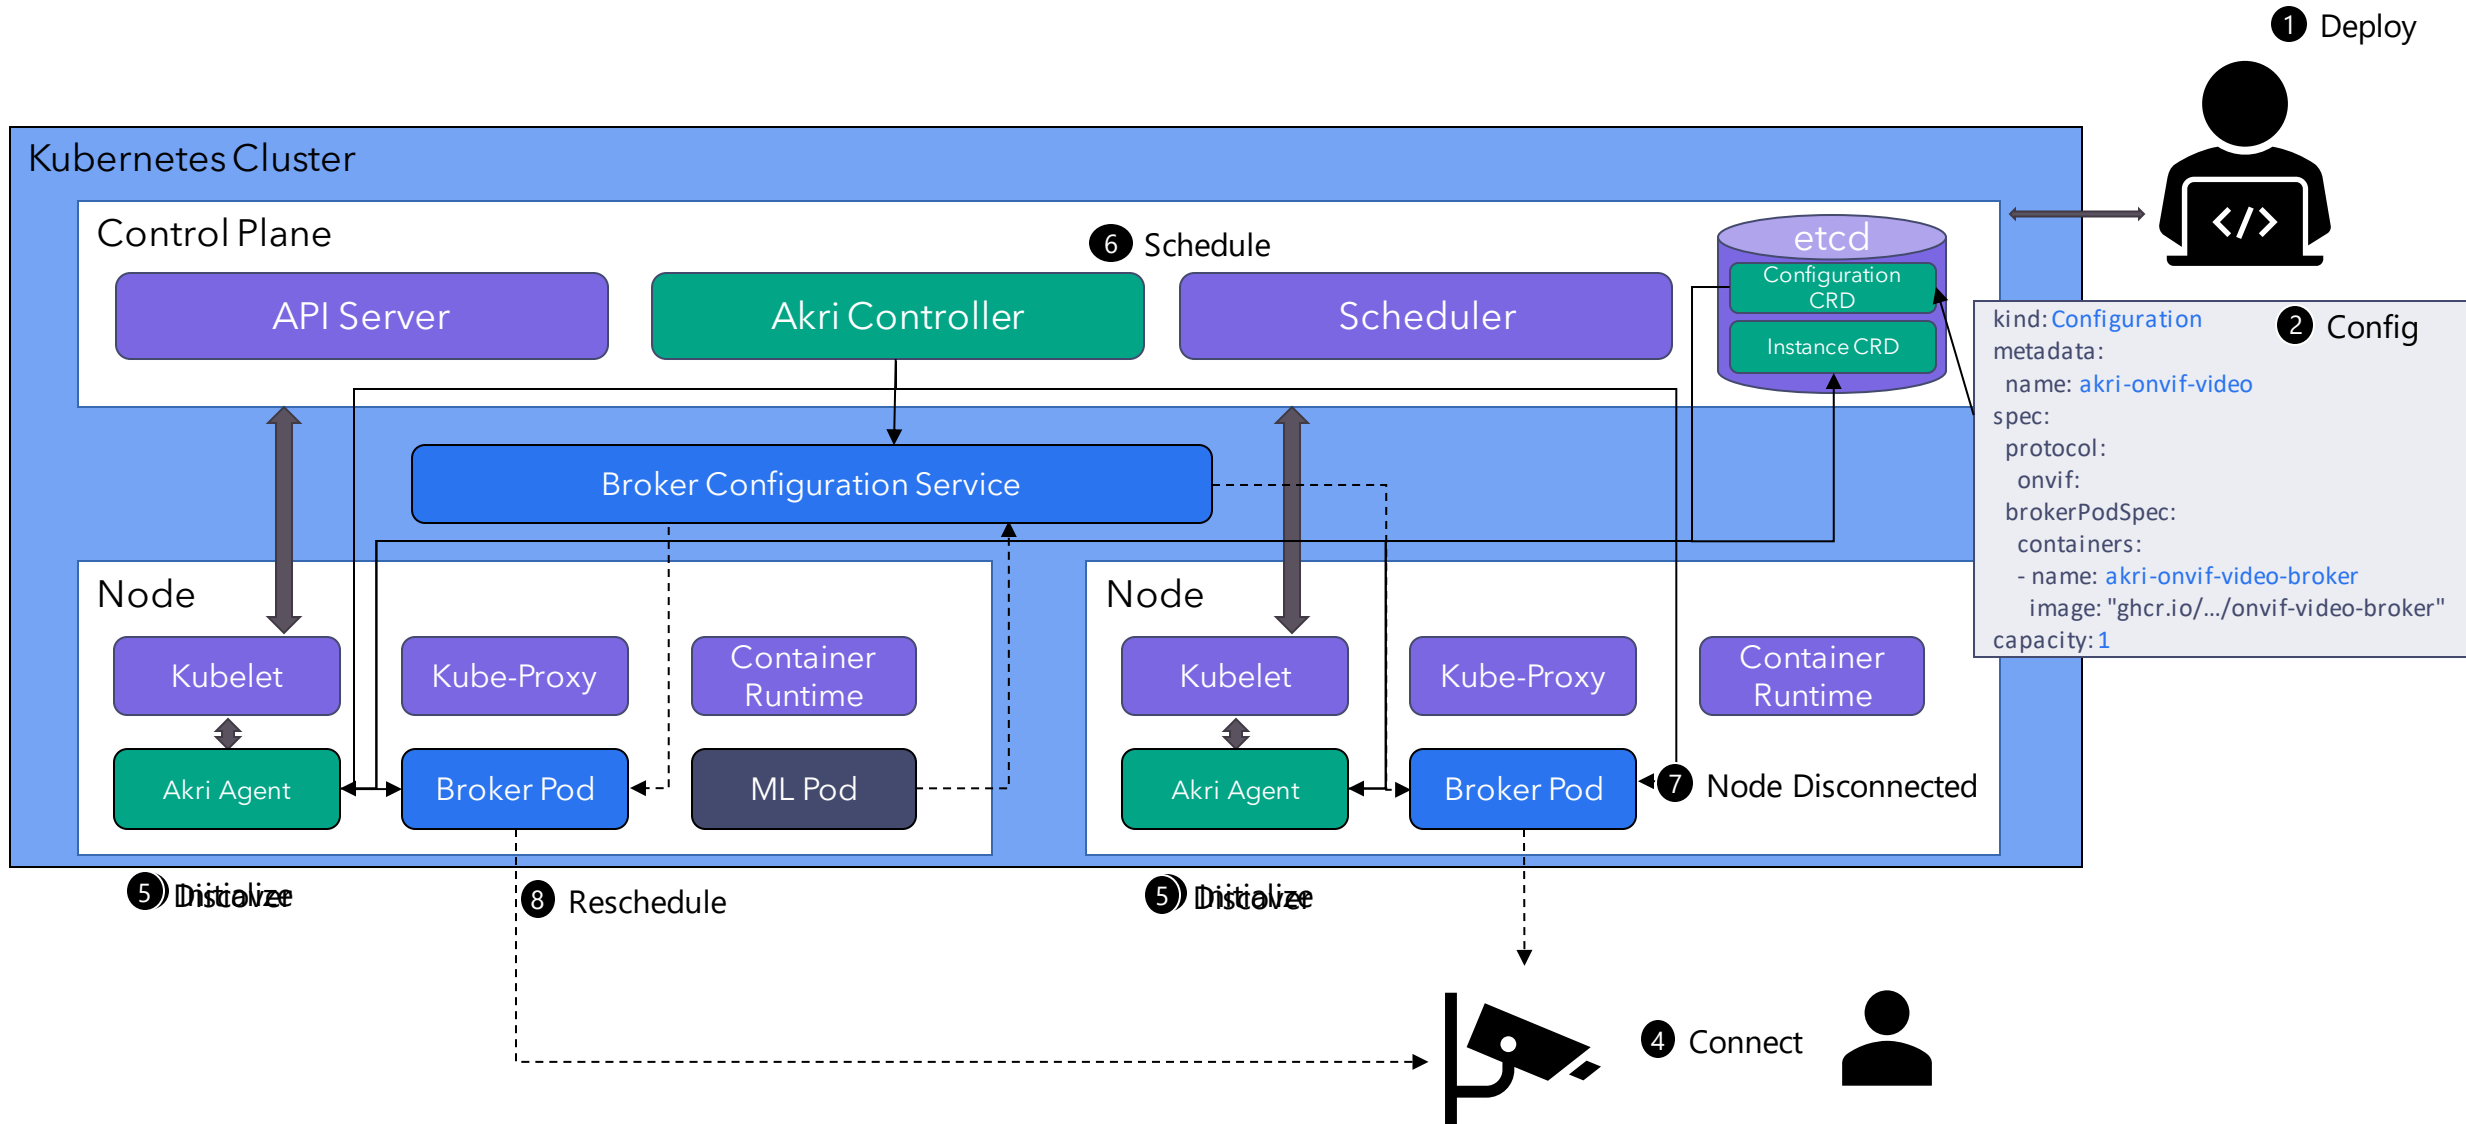
Akri Architecture



Kubernetes Device Plugin Refresher



Akri Flow



Akri Functionality



Resource Interface

Abstraction for declaring what devices should be discovered

Device Discovery

Discovery of leaf devices and graceful handling of intermittent connectivity and removals

Broker Deployment

Broker deployment strategy similar to DaemonSet

High Device Availability

HA of leaf devices via "sharing" / capacity

Automated Services

Targeting specific devices or device type

Support for Multiple Protocols

Currently supported:
ONVIF and udev

Extensible

BYO protocol handlers and broker pods

Standards Based

Application specific protocol translation gateways

Akri's Four Components

Working together so “You name it, Akri finds it, You use it”



Configuration CRD –
where “you name it”



Akri Agent – how
“Akri finds it”



Instance CRD – a
representation of “it”



Akri Controller –
helps “you use it”

Demo

Roadmap



Additional discovery protocols and improving extensibility



Additional deployment strategies



Improving composability

Most importantly - We want to hear from you!

Together, we can electrify the edge for Kubernetes users

1

Try it out

End to end demo
discovering mock
cameras

aka.ms/akri/e2e

2

Learn more

Akri docs

aka.ms/akri/docs

3

Tell us what you think

GitHub Issues
Akri channel

aka.ms/akri/channel

4

Help us define the future for Akri

Proposals


aka.ms/akri/proposals

5

Help us discover more

Contribute

aka.ms/akri/contributing



github.com/deislabs/akri